

SIVANAGARAJU BORRAJU

sivaborraju@gmail.com · (409) 867-6900 · [linkedin.com/in/sivanagaraju09](https://www.linkedin.com/in/sivanagaraju09)

PROFESSIONAL SUMMARY

Data Engineer with 5+ years of experience building enterprise data pipelines, dimensional models, and analytical solutions using Python, SQL, PySpark, Apache Spark, Azure Databricks, Azure Data Factory, Snowflake, and Airflow. Experienced in data warehouse and data mart design, batch and streaming processing, data quality, governance, CI/CD, and production troubleshooting. Built curated Snowflake datasets supporting Power BI reporting and enterprise analytics.

TECHNICAL SKILLS

- **Programming & Scripting:** Python (pandas, numpy, scikit-learn), SQL (joins, CTEs, window functions, partitioning, clustering, and stored procedures), PySpark, Spark SQL, Java (OOP), Shell Scripting (UNIX/Linux)
- **Big Data & ETL/ ELT Frameworks:** Apache Spark (Core, SQL, Structured Streaming), Hadoop M/R, Apache Hive, Delta Lake, Data Lake
- **Cloud Platforms:** Azure (Databricks, Azure Data Factory, Data Lake Storage Gen2, Blob Storage, Synapse Analytics, Unity Catalog), AWS (EMR, EC2, S3, Redshift), GCP (Google BigQuery, GCS, Dataflow)
- **Workflow Orchestration & Streaming:** Databricks Workflows, Apache Airflow, Apache Kafka
- **Data Warehousing & Modeling:** Snowflake, Logical & Physical Data Modeling, Dimensional/Star Schema, Fact & Dimension Tables, Data Marts, Relational Table Design (PK/FK)
- **Databases:** Relational (MS SQL Server, MySQL, PostgreSQL), No-SQL (MongoDB, Cassandra)
- **Visualization & Reporting:** Power BI, Tableau
- **DevOps & CI/CD:** Git, GitHub, Jenkins, Jira, Confluence, Version control, access control
- **AI & Methodologies:** Agile (Scrum), Data Governance & Data Quality Frameworks, GitHub Copilot

PROFESSIONAL EXPERIENCE

Data Engineer | AT&T

August 2023 – Present

Dallas, TX

- Architected ETL/ELT data pipelines using Azure Databricks, PySpark, Spark SQL, and Azure Data Factory, transforming structured and semi-structured data into curated datasets for the Snowflake data warehouse.
- Built Kafka and Spark Structured Streaming pipelines processing 50K events/sec and 5TB of CRM, billing, and payment data, enabling near-real-time analytics and improving job performance by 35%.
- Developed logical and physical data models, defining business entities, fact and dimension structures, relational table design, and data grain to support reusable data warehouse and data mart designs.
- Built subject-area data marts in Snowflake from curated dimensional datasets, structuring business-ready data for reporting and Power BI dashboards.
- Engineered bronze, silver, and gold data layers using Delta Lake and Medallion Architecture, applying standardized transformation rules to improve consistency and downstream reliability.
- Optimized Spark applications through partitioning, broadcast joins, caching, and Adaptive Query Execution, reducing production processing time by 40%.
- Implemented data quality checks, reconciliation, schema validation, duplicate/integrity detection, performance optimization, and reduced downstream reporting issues by 30%.
- Applied Unity Catalog for data access governance, supporting data governance documentation, dataset ownership, metadata, and lineage practices for downstream consumers.
- Automated job scheduling and orchestration with Airflow and Databricks Workflows, managing dependencies, monitoring, and failure recovery to maintain 99%+ SLA compliance.
- Used GitHub Copilot for SQL development, debugging, detect pipeline failures, testing, deployment, and technical documentation, validating outputs for correctness and data sensitivity before production.

- Supported CI/CD deployments (Git, Jenkins) and production troubleshooting in Agile delivery, collaborating with product owners, QA, and platform engineers, delivered curated Snowflake datasets for Power BI reporting to business stakeholders.

Tools Used: Azure Databricks, Apache Spark, PySpark, Spark SQL, Delta Lake, Spark Structured Streaming, Apache Kafka, Azure Data Factory, Azure Data Lake Storage Gen2, Azure Synapse Analytics, MS SQL Server, Snowflake, Unity Catalog, Databricks Workflows, Airflow, Python, SQL, JSON, Parquet, REST APIs, Git, Jenkins, CI/CD, Power BI, Tableau

Data Engineer | Infosys

April 2020 – December 2022

India

- Built ETL/ELT pipelines using PySpark, Python, SQL, Azure Data Factory, and Snowflake, transforming large structured and semi-structured datasets into curated data for BI and reporting.
- Designed dimensional data models using fact and dimension tables and relational table structures (PK/FK) to support enterprise reporting and analytical workloads.
- Integrated REST APIs, JSON, flat files, and relational database sources (MySQL, PostgreSQL via JDBC) into Azure Data Lake Storage and Snowflake, establishing reliable data flows for downstream analytics.
- Implemented reconciliation, schema validation, incremental loading, and automated data quality checks, improving production data reliability.
- Developed reusable PySpark and SQL transformation frameworks, standardizing processing patterns and reducing duplicated engineering effort.
- Automated job scheduling and orchestration with Airflow, managing task dependencies, monitoring, retries, and failure handling for recurring pipelines.
- Modernized legacy pipelines by migrating workloads to Azure Databricks and Azure Data Lake Storage while maintaining production stability.
- Automated CI/CD deployment (Git, Jenkins) for Python, SQL, and Spark applications within Agile delivery teams, supporting consistent promotion across environments.

Tools Used: Azure Databricks, Apache Spark, PySpark, SparkSQL, Python, SQL, Azure Data Factory, Azure Data Lake Storage Gen2, Azure Synapse Analytics, Snowflake, Airflow, JSON, Parquet, REST APIs, Git, Jenkins, CI/CD

CERTIFICATIONS

Microsoft Certified: Fabric Data Engineer Associate (DP-700)

EDUCATION

Master's Degree in Computer Science | Lamar University, Beaumont, TX

Bachelor's Degree in Computer Science | Vignan's Lara Institute of Technology & Science, India